

From Understanding to Execution: The R-CoT Method for Structuring Reasoning in Large Language Models

*(An empirical evaluation across **10**
models and **56** tasks; R-CoT success
rate = **(96.4%)**)*

Omar Abdallah

*Creator, Developer, and Evaluator of the Reflective
Chain-of-Thought (R-CoT) Technique*

20091512o@gmail.com

Abstract

We present Reflective Chain-of-Thought (R-CoT), a prompting framework composed of three sequential stages:

*(1) **Task Understanding** – rephrasing the task and stating any necessary assumptions;*

*(2) **Reflective Reasoning** – generating a numbered plan of 3–6 steps to structure the solution process;*

(3) **Execution** – producing the final answer based on the reasoning plan.

We compared R-CoT with two baseline prompting strategies—Direct Prompting and Chain-of-Thought (CoT)—through a structured large-scale experimental study.

Models tested: GPT-4o, GPT-4, GPT-3.5 Turbo, Claude Sonnet 4, Gemini Flash 2.5, Mistral 7B, Command R, LLaMA 3.3 70B, Grok-3, and DeepSeek V3.

Tasks: 56 tasks per technique (168 total evaluations), spanning 19 distinct task categories, including writing, summarization, code debugging, programming, analysis, brainstorming, information structuring, critical thinking, persuasion, deep reflection, prompt engineering, explanation, rewriting, code explanation, creative ideation, educational design, visual structuring, instructional critique, pedagogical reasoning, instructional simplification, UX critique, and rational thinking.

Evaluation: One human rater (the author) using both manual assessment and AI-assisted tools.


Statistical analysis: Main test — Friedman test ($\chi^2 \approx 36.11$, $p < 0.000000015$), showing a highly significant difference between the three techniques; Post-hoc — Wilcoxon signed-rank test with Bonferroni correction:

- R-CoT vs CoT: $p \approx 0.000011$ (strongly significant in favor of R-CoT)
- R-CoT vs Direct: $p \approx 0.000011$ (strongly significant in favor of R-CoT)
- CoT vs Direct: $p \approx 0.00049$ (significant in favor of CoT)

Key results: R-CoT achieved 54 wins out of 56 rounds (success rate = 96.4%), compared to 2 wins for CoT and 0 wins for Direct Prompting.

Conclusion: The results demonstrate that incorporating an explicit reflective stage into the prompting process (R-CoT) substantially improves performance across a wide variety of cognitive tasks, achieving a 96.4% success rate and outperforming both CoT and Direct Prompting by a significant margin. This establishes R-CoT as a generalizable and effective prompting strategy for enhancing reasoning in large language models. Moreover, these findings suggest that reflective prompting not only improves current output quality but also opens promising pathways for developing more transparent, reliable, and interpretable AI systems in the future.

Figure 1. The three stages of the Reflective Chain-of-Thought (R-CoT) technique.



Introduction & Related Work

In recent years, Large Language Models (LLMs) have demonstrated remarkable progress in their ability to perform a wide variety of tasks, ranging from translation and text summarization to programming and statistical analysis. However, their ability to carry out structured reasoning and solve multi-step problems remains a significant challenge, especially in scenarios that require a precise logical sequence to arrive at a correct answer.


The Chain of Thought (CoT) methodology emerged as a prominent approach to address this challenge by prompting the model to explicitly generate intermediate reasoning steps. While CoT has achieved notable improvements in reasoning accuracy in certain domains, its application in diverse, multi-task settings has revealed limitations in generalization and consistency. CoT-generated reasoning chains may include unnecessary or redundant steps, leading to inaccurate or contradictory conclusions (see Table 1).

Table 1 – Key differences between prominent prompting techniques

Technique	Core Mechanism	Primary Benefit	Limitation / Focus
Direct Prompting	Single instruction, immediate response	Simplicity, speed	Lacks reasoning depth, prone to errors

Chain of Thought (CoT)	Step-by-step reasoning	Improved reasoning on complex tasks	May lack initial understanding or structured planning
Self-Consistency	Multiple CoT paths, majority vote	Robustness, error reduction	Computationally intensive, no explicit initial reflection
ReAct	Reasoning + external actions	Dynamic interaction, tool use	Relies on external tools, potential latency
Tree of Thought (ToT)	Tree-like exploration of reasoning paths	Deliberate planning, complex problem-solving	Computationally demanding, initial understanding not enforced
Reflective CoT (R-CoT)	Three-step reflective process (Understand → Reason → Act)	Deep understanding, structured planning, enhanced output quality	Requires explicit prompting for each step

To clarify the structural difference between CoT and R-CoT, **Figure 2** illustrates a side-by-side process diagram for both methods.



The development of R-CoT was motivated by the need to bridge the gap between strictly structured reasoning and maintaining flexibility and efficiency. Unlike traditional CoT, R-CoT integrates structured reasoning steps with a reflective component, enabling the model to review and adjust its reasoning pathway during generation. This mechanism reduces errors caused by early deviations in

reasoning and improves the final output's accuracy across multiple task types.

A practical consideration is that the current form of R-CoT requires explicit prompting for each stage (i.e., separate instructions for Understand, Reflective Reasoning, and Act), which may increase input length or the number of interactions with the model. To address this, we propose the development of a hybrid R-CoT version that merges the three stages into a single, unified message prompt while internally preserving the step division (i.e., embedded understanding, structured plan, and reflective check before execution). The design, examples, and evaluation framework for this single-message hybrid R-CoT will be discussed in Section X – Future Work.

To illustrate the practical impact of this methodology, **Table 2** presents verbatim excerpts from a literary analysis task (task 5), showing the actual outputs produced under Direct Prompting, CoT, and R-CoT, together with a brief neutral observation for each.

Table 2 – Verbatim excerpts from a literary analysis task (Day 5)

Technique	Verbatim output (excerpt)	Neutral observation
Direct Prompting (Normal result)	<i>This excerpt—"And the moon leaned closer to the river, whispering secrets only the waves could understand..."—is rich with lyrical symbolism and evocative imagery. The moon traditionally symbolizes mystery, femininity, and distant wisdom, while the river represents emotional flow, time, and subconscious depth. When the moon "leans closer," it personifies celestial intimacy, as if the distant observer seeks to connect with the ever-changing river, perhaps echoing a desire for communion between the eternal and the ephemeral. The phrase "whispering secrets" implies a quiet, almost sacred exchange—intimate and invisible to the outside world. The river's waves, capable of understanding these whispers, suggest that nature holds a language beyond human comprehension, one born of rhythm and intuition. Metaphorically, this could reflect the relationship between thought (the moon) and feeling (the river), or between the rational and the intuitive self. The emotional tone is gentle and wistful, imbued with a sense of longing and quiet connection. Aesthetically, the passage reads like a soft nocturne—a brief moment of poetic harmony between sky and earth, thought and feeling, silence and motion. In my interpretation, the excerpt suggests that some truths or desires can only be understood in stillness and through deep emotional listening.</i>	A concise, interpretive summary capturing imagery and tonal features; presents a final interpretation without exposing intermediate analytical steps.

Chain-of-Thought (CoT)	<i>Let's walk through the Chain of Thought (CoT) step-by-step to analyze the excerpt: Excerpt: "And the moon leaned closer to the river, whispering secrets only the waves could understand..." --- 1. Identify the Type of Text This is clearly poetic and metaphorical. The sentence uses personification and lyrical language, suggesting it comes from a poem or poetic prose. It's not literal or narrative—it evokes a symbolic interaction between elements of nature. --- 2. Look for Imagery, Metaphors, or Symbols "The moon leaned closer" – personifies the moon, giving it intent and movement. "Whispering secrets" – suggests intimacy and mystery. "Only the waves could understand" – gives consciousness to the river, implying a private communication between natural elements. This is rich in imagery and relies on symbolic personification of nature to evoke a deeper emotional landscape. --- 3. Analyze the Emotional Tone and Mood The tone is gentle, mysterious, and intimate. Words like "leaned," "whispering," and "secrets" evoke a quiet moment of closeness and trust. The mood feels dreamlike and romantic, perhaps even melancholic—a nighttime solitude where hidden emotions or unspoken truths emerge. --- 4. Interpret the Symbolic Meaning Moon: Often symbolizes reflection, the unconscious, emotion, or feminine energy. River/Waves: Symbolize life's flow, emotions, or memory—especially when moving or responsive like waves. The act of the moon whispering to the river could symbolize a moment of emotional revelation, intuition sharing truths with the emotional self. The secrecy implies that some truths are only understood by those attuned to deeper, non-verbal, or natural rhythms—like the waves. Symbolically, this line may suggest the idea of hidden wisdom or emotional truths being communicated quietly within oneself, where only the emotional or intuitive part can truly grasp them. --- 5. Consider the Form or Rhythm Though just one sentence, the rhythmic phrasing mimics a poetic line: Balanced clauses: "And the moon leaned closer / to the river, whispering secrets / only the waves could understand." It flows like a soft wave, mirroring the content itself. This musicality reinforces the mood of calm mystery and night-time serenity. --- 6. Final Interpretation This poetic line uses personification and natural imagery to evoke a secret, emotional exchange between celestial and earthly forces. The moon, as a symbol of introspection and silent wisdom, shares its quiet truths with the ever-moving river,</i>	A detailed, step-by-step analytical trace showing the CoT procedure (type → imagery → tone → symbolism → rhythm → interpretation); exposes intermediate reasoning but lacks an explicit reflective verification stage before the final output.
--------------------------------------	---	---

	<i>which represents emotion and the unconscious. The beauty of this moment lies in its quietness—these are not truths for the world, but for those attuned to deeper, internal currents. Emotionally, it conveys intimacy, solitude, and the unseen bonds between forces in nature—and perhaps within the self.</i>	
--	--	--

Reflective CoT (R-CoT)	<i>Absolutely! Let's begin with Step 1 – Understand. --- ✓</i> <i>Step 1 – Understand From the task, I understand that you want a symbolic or aesthetic analysis of the given poetic excerpt. That means my role is not just to paraphrase or explain literally, but to uncover: Symbolic meanings: What deeper ideas or archetypes the moon, river, and waves might represent. Emotional tone: What kind of mood or atmosphere is created. Artistic techniques: Any use of metaphor, personification, or poetic devices. The Given Excerpt: "And the moon leaned closer to the river, whispering secrets only the waves could understand..." ---</i> <i>My Planned Approach for Step 2: 1. Break down the sentence into symbolic components: What the moon might symbolize. What the river and waves could symbolize. 2. Analyze the interactions (e.g., moon whispering, river listening). 3. Identify poetic techniques (e.g., personification, alliteration, mood). 4. Reflect on the emotional and thematic atmosphere. 5. Consider if there's a metaphysical, romantic, or existential subtext. --- Go to step two ... ✓</i> <i>Step 2 – Think (symbolic breakdown of Moon, River, Waves, 'whispering secrets', artistic effects, metaphysical layer) ... Go to step 3 ... ✓</i> <i>Step 3 – Act In the poetic line, "And the moon leaned closer to the river, whispering secrets only the waves could understand," the imagery evokes a rich, dreamlike atmosphere filled with intimacy and quiet mystery. The moon, often a symbol of intuition, feminine energy, and the hidden self, is brought to life through personification—it "leans" and "whispers," breaking its usual role as a passive observer. Its gesture suggests a rare closeness, as if divine or cosmic wisdom is descending toward the emotional world of the river. The river, a classical symbol of life's journey, emotional depth, and the passage of time, becomes the quiet listener to this celestial confession. Yet, it is not the river itself that understands, but the waves—subtle, shifting, and sensitive—who register and interpret the whispered truths. This detail implies that true understanding lies not in the broad current of life, but in the emotional ripples, the quiet reactions we often overlook. The line subtly captures the elusive way in which nature communicates with itself—and with us—suggesting that some truths are too delicate for words, audible only to those attuned to silence and movement. Overall, the excerpt weaves an aesthetic tapestry of melancholy intimacy, cosmic emotion, and the beauty of unspoken understanding.</i>	Shows an explicit three-stage structure (Understand → Think → Act) with planned steps and a final integrated output; provides the highest transparency in reasoning flow and includes internal planning and verification.
--------------------------------------	--	--

Provenance note: The excerpts in Table 2 are **verbatim outputs** from the experimental runs (Day

5). Full prompts and complete outputs for all trials are provided in **Appendix A** to enable replication and verification.

Compared to previous prompting techniques, R-CoT demonstrates greater consistency and performance, with enhanced adaptability to the nature of the task—whether analytical, programming-related, or linguistic. These contributions warrant a comprehensive experimental evaluation, which we detail in the following sections.

The Method

The Reflective Chain-of-Thought (R-CoT) technique is designed to enhance the reasoning capabilities of Large Language Models (LLMs) through a structured three-stage process: Understand, Reason, and Act.

This approach ensures that the model not only produces the final answer but also verifies and refines its reasoning pathway before generating the output.

1. Understand

In this stage, the model carefully analyzes the task, identifies requirements, and clarifies any implicit constraints.

Instead of proceeding directly to the solution, the model reformulates the task in its own words to ensure complete comprehension.

***Example:* When asked to summarize the contributions of a paper in three bullet points, the model first identifies the paper’s topic, determines the required output format, and focuses on the contributions rather than methodological details.**

2. Reason

The model develops a structured logical pathway for solving the task, breaking it down into coherent sub-steps.

The objective is to create a clear and verifiable reasoning process.

Example: Read the abstract and conclusion to extract contributions, cross-check with the methodology to ensure completeness, and prioritize clarity and conciseness.

3. Act

The model generates the final output based on the organized reasoning developed in the previous stage.

This phase also includes a self-review to verify that the output meets all task requirements.

Example: Produce the three required bullet points with precision, ensuring compliance with all constraints.

Rationale for the Step Sequence

- **Understand:** Prevents early errors caused by misinterpretation of the task.
 - **Reason:** Encourages organized and transparent thinking.
 - **Act:** Ensures high-quality outputs that meet all requirements.
-

Experimental Environment and Design

The R-CoT technique was evaluated against two baseline methods:

Chain-of-Thought (CoT) and Direct Prompting, across 10 LLMs:

GPT-4o, GPT-4, GPT-3.5 Turbo, Claude Sonnet 4, Gemini Flash 2.5, Mistral 7B, Command R, LLaMA 3.3 70B, Grok-3, and DeepSeek V3.


- **Tasks:** 19 distinct types, including writing, summarization, code debugging, programming, analysis, brainstorming, information structuring, critical thinking, persuasion, deep reflection, prompt engineering, explanation, rewriting, code explanation, creative ideation, educational design, visual structuring, instructional critique, pedagogical reasoning, instructional simplification, UX critique, and rational thinking.

- **Task Distribution:** Randomly assigned across models.
- **Settings:** All models were used through their official applications with default settings; the primary execution language was English.
- **Applied Techniques:** R-CoT, CoT, Direct Prompting.
- **Evaluation:** Each model's outputs were assessed using four criteria — clarity, structure, depth, and creativity. For each task, results from the three techniques were compared, and the best-performing approach was selected.

The evaluation encompassed a total of 168 runs (56 tasks × 3 prompting techniques), ensuring a robust and balanced comparison. Presenting the experimental design (Figure 3) serves not only to describe the setup but also to promote transparency and enable replication of the study.

Figure 3 presents the experimental design, illustrating the flow from models to tasks, then to techniques, followed by evaluation and the selection of the best technique for each task-model pair.

Figure 3: Experimental design for evaluating R-CoT against CoT and Direct Prompting.



Note: A full step-by-step example of R-CoT in practice is presented in the Introduction & Related Work section, enabling readers to trace the methodology from instruction to final output.

Introduction & Related Work section.


Results & Discussion

The **Reflective Chain-of-Thought (R-CoT)** technique was evaluated against two baseline methods — **Chain-of-Thought (CoT)** and **Direct Prompting** — across **56 tasks** spanning 19 distinct categories.

In total, the evaluation comprised **168 runs** (56 tasks × 3 prompting techniques), with each task category assessed independently to determine the winning technique based on performance criteria.

As shown in **Figure 4 (Number of wins for each technique across all experimental runs — R-CoT, CoT, Direct Prompting)**

First-place wins per technique (out of 56 evaluations)



R-CoT achieved a dominant **54 wins out of 56** (success rate = 96.4%), compared to only **2 wins** for CoT and **0 wins** for Direct Prompting.

This substantial performance gap underscores the consistent advantage of integrating a reflective reasoning stage into the problem-solving process.

Beyond the number of wins, **R-CoT** achieved an average task accuracy of **96.4%**, outperforming both CoT and Direct Prompting in nearly all task types. Statistical analyses confirmed the robustness of these differences: the **Friedman test** yielded $\chi^2 \approx 36.11$ with **$p < 0.000000015$** (highly significant), and post-hoc **Wilcoxon signed-rank tests** with Bonferroni correction showed:

- **R-CoT vs CoT:** $p \approx 0.000011$ (strongly significant in favor of R-CoT)
- **R-CoT vs Direct:** $p \approx 0.000011$ (strongly significant in favor of R-CoT)
- **CoT vs Direct:** $p \approx 0.00049$ (significant in favor of CoT)

When analyzed by task type, **R-CoT** maintained high accuracy with low variance in analytical reasoning tasks, whereas **CoT** exhibited greater variability and occasional reasoning drift. In programming and code-debugging tasks, the **Understand** stage in

R-CoT reduced syntax and logical errors, resulting in higher execution success rates.

From a qualitative perspective, R-CoT outputs were not only more accurate but also more structured and interpretable, enabling human evaluators to verify correctness more efficiently. This interpretability is particularly valuable for applications requiring high reliability and transparency, such as educational content generation, decision support, and complex analytical problem-solving.

Overall, the findings indicate that the primary performance advantage of R-CoT stems from its clear separation of the **Understand** → **Reason** → **Act** stages. In contrast, CoT often blended these stages, making it more susceptible to early reasoning errors that propagated into the final output.

Conclusion & Future Work

This study introduced and applied the **Reflective Chain-of-Thought (R-CoT)** technique as an improvement over traditional **Direct Prompting** and **Chain-of-Thought (CoT)** methods.

By incorporating a reflective stage prior to reasoning, R-CoT enables Large Language Models (LLMs) to:

1. **Understand the task and its requirements** with greater precision.
2. **Plan a coherent reasoning path** before execution.
3. **Produce outputs that are more structured, accurate, and context-aware.**

Practical comparisons demonstrated that **R-CoT** delivers more complete and balanced responses, particularly for complex or multi-dimensional tasks where premature reasoning may lead to errors or omissions.

The overall success rate reached **96.4%** (54 wins out of 56 tasks) with strong statistical significance, underscoring the method's effectiveness in enhancing output reliability and interpretability.

The significance of R-CoT lies in its resemblance to expert problem-solving processes

— **reflecting and planning before acting** — which enhances both transparency and trustworthiness in generated results.

Future Work

Building on the current results, future research directions may include:

- **Fast variant for simple tasks:** Develop a streamlined version of R-CoT that merges the Understand, Reason, and Act stages into a single prompt, while preserving the internal reasoning structure, to reduce processing time for low-complexity tasks.
- **Integration with more complex techniques for larger tasks:** Explore combining R-CoT with advanced methods such as **Tree of Thought (ToT)** to address large-scale and complex problems, leveraging deep planning capabilities while retaining the clarity of the reflective stages.

Figure 5 (Timeline of Future Work) presents the proposed development roadmap for these future research directions.



Final Remarks

The **R-CoT** technique has the potential to enhance prompt engineering practices by

combining simplicity with effectiveness.

As LLMs continue to evolve, adapting R-CoT for both fast and complex tasks could pave the way for more flexible, reliable, and transparent AI systems, benefiting both academic research and real-world applications.

References

- [1] J. Wei, X. Wang, D. Schuurmans, M. Bosma, B. Ichter, F. Xia, E. H. Chi, Q. V. Le, and D. Zhou, "Chain-of-Thought Prompting Elicits Reasoning in Large Language Models," *Advances in Neural Information Processing Systems*, vol. 35, pp. 24824–24837, 2022.
- [2] X. Wang, J. Wei, D. Schuurmans, Q. V. Le, E. H. Chi, S. Narang, A. Chowdhery, and D. Zhou, "Self-Consistency Improves Chain-of-Thought Reasoning in Language Models," in *Proceedings of the International Conference on Learning Representations (ICLR)*, 2023.
- [3] S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao, "ReAct: Synergizing Reasoning and Acting in Language Models," *arXiv preprint arXiv:2210.03629*, 2022.
- [4] S. Yao, D. Yu, J. Zhao, I. Shafran, T. L. Griffiths, Y. Cao, and K. Narasimhan, "Tree of Thoughts: Deliberate Problem Solving with Large Language Models," *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [5] J. Zhou, N. Zhang, X. Ren, and M. Y. Kan, "Least-to-Most Prompting Enables Complex Reasoning in Large Language Models," *arXiv preprint arXiv:2205.10625*, 2022.
- [6] Z. Liu, J. Li, H. Zhang, C. Wang, and K. Xu, "Automatic Chain of Thought Prompting in Large Language Models," *arXiv preprint arXiv:2302.13845*, 2023.
- [7] L. Gao, M. Zhao, and T. Yu, "PAL: Program-Aided Language Models," *arXiv preprint arXiv:2211.10435*, 2023.

Appendix

A. Experimental Results

The full experimental runs, metrics, and logs are provided in the supplementary PDF.

Link to Experimental Results PDF:

[[Full experiment](#)]

B. Optimal Settings

The detailed task-settings mapping (winning configurations) is provided in the settings document.

Link to Settings PDF: [[setting experiment](#)]

C. Source Code

The final implementation of the Reflective Chain-of-Thought (R-CoT) framework, including task classification, heuristics, and automated prompt generation, is available as a direct downloadable file.

Download the code here: [[link of the code](#)]

D. GitHub Repository

The full repository with version control, issue tracking, and potential future updates is hosted on GitHub.

Visit the repository here: [[Link of GitHub Repository](#)]

E. license

> © 2025 Omar. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License (CC BY-NC-SA 4.0).

> Full license text: [LICENSE R-CoT.pdf](#)